

Main system comparison

Dots show mean feasible utility; bars show 95% bootstrap intervals over cards. LLM rows show final output after validation/repair.

