

Corrected v0.1.0 benchmark pipeline

One fixed action contract • scorer outcomes never enter deployable system prompts • raw and repaired views share the same responses

CARA v1.0.1 LO_All → corrected positional import and exhaustive audit

100 official tasks • include every task with at least 10 feasible candidates → 91 included / 9 excluded

91 SYSTEM-INPUT CARDS

assay context • 50 measured supports • finite candidate pool • constraints • k=10
candidate activity absent

91 SCORER-ONLY ROWS

candidate pChEMBL outcomes
same task_id + system-input SHA256

7 DETERMINISTIC SYSTEMS

6 deployable: random-valid, rules, similarity,
QSAR RF / GBT / linear SVR
1 scorer-only oracle upper bound

6 RAW LLM CONDITIONS

3 frozen provider/model conditions × 2 interfaces
91 cards per condition = 546 provider requests
exact requests • one successful attempt • cached traces

6 ZERO-CALL REPAIRED VIEWS

91 actions per view
retain valid IDs • remove violations
fill from deterministic rules ranking

evaluator/oracle only

same response

ORDERED TOP-10 ACTION RECORDS

7 deterministic/oracle traces • 6 raw LLM traces × 91 • 6 post-hoc-repaired traces × 91 • repair adds zero provider calls

SCIENTIFIC UTILITY / RANKING

feasible utility • NDCG@10
constrained regret • oracle headroom
selected IDs joined to scorer outcomes

ACTION VALIDITY / FAILURES

zero-issue whole-action validity
valid-selection compliance • schema / contract /
constraint failure taxonomy

PAIRED 91-CARD ANALYSIS

task-level 95% bootstrap intervals
QSAR vs repaired/raw LLM • bare vs descriptors
raw vs repaired attribution

CANONICAL 19-SYSTEM COMPARISON

primary leaderboard • paired deltas • card diagnostics • raw/final ablations • latency, tokens, usage-derived cost