

Raw LLM failure taxonomy (91 cards per condition)

Anthropic claude-opus-4-8 bare	17 (18.7%)	18 (19.8%)	44 (48.4%)	46 (50.5%)
Anthropic claude-opus-4-8 descriptors	22 (24.2%)	7 (7.7%)	42 (46.2%)	45 (49.5%)
DeepSeek deepseek-v4-pro bare	8 (8.8%)	22 (24.2%)	51 (56.0%)	46 (50.5%)
DeepSeek deepseek-v4-pro descriptors	22 (24.2%)	5 (5.5%)	32 (35.2%)	46 (50.5%)
OpenAI gpt-5.5-2026-04-23 bare	72 (79.1%)	13 (14.3%)	0 (0.0%)	10 (11.0%)
OpenAI gpt-5.5-2026-04-23 descriptors	68 (74.7%)	14 (15.4%)	0 (0.0%)	10 (11.0%)
	No detected issue	JSON/schema failure	Selection-contract failure	Candidate-constraint failure

Green denotes zero-issue actions. Failure categories can overlap, so row counts must not be summed.